

# TERRAFORM

**laC:** Infrastructure as Code. Write and execute code to define, deploy, update, and destroy your infrastructure.

## **laC Categories:**

- **Ad hoc scripts:** Manual scripts which developers write in their favorite language to perform an ad-hoc (one time) task.
- **Configuration Management Tools:** Ex: Chef, Ansible, Puppet etc. Provides the ability to install and manage software on existing resources. Declarative in nature (How you want?)
- **Server Templating tools:** Ex: Docker, Packer, Vagrant. Instead of launching server(s), create an image of the server (snapshot) which is self-contained and can be deployed. Guarantees that the image works the same in all environments (dev, staging and prod).
- **Orchestration tools:** Ex: Kubernetes, ECS, Docker Swarm etc. Used for managing the server templating tools.
- **Provisioning tools:** Ex: Terraform, CloudFormation, Resource Manager, Cloud Deployment. Allows for creation of resources. Declarative in nature (What you want?)

## **Benefits of laC:**

- **Self-service**
- **Speed and safety**
- **Documentation and Version control**
- **Validation**
- **Reuse**
- **Idempotence:** Running the code which produces same result when ran several times is called Idempotence

**Flow:** Terraform configuration file -> Provider -> Cloud

**Provider:** Cloud provider - AWS, Google, Azure etc. A provider is responsible for understanding API interactions and exposing resources.

```
provider "aws" {  
    region    = "us-west-2"  
    access_key = "my-access-key" # Use only for development  
    secret_key = "my-secret-key" # Use only for development  
  
    # assumes role to perform action instead of access and secret access keys  
    assume_role {
```

```

        <details of role to assume>
    }
}

```

- Another way is to use aws config

- "alias" will allow to create two same provider configurations

**Example:**

```

provider "aws" {
    region = "ap-southeast-1"
}
provider "aws" {
    region = "ap-southeast-2"
    alias = "sydney"
    Profile = "<profile name>" # here profile name refers to the AWS profile
                                created using aws cli
}

resource "aws_instance" "MyEC1" {
    instance_type = "t2.micro"
    provider = "aws.sydney"
}

```

**Provider version notations:**

- >=1.0: Greater than equal to version 1.0
- <=1.0: Less than equal to version 1.0
- ~>2.0: Any version in 2.x range
- >=2.0,<=2.30: Any version between 2.0 and 2.30

**Configuration files:** Declarative representation of the infra. Should end with .tf or .tf.json

**Resources:**

- Describes one or more infrastructure objects such as VM, LB etc

```

resource "aws_instance" "web" {
    ami      = "ami-a1b2c3d4"
    instance_type = "t2.micro"
}

```

**Important arguments:**

**Depends on:** To handle hidden resource dependencies that Terraform can't automatically infer.

Example: [https://www.terraform.io/docs/configuration/resources.html#depends\\_on-explicit-resource-dependencies](https://www.terraform.io/docs/configuration/resources.html#depends_on-explicit-resource-dependencies)

Lifecycle: Affects how terraform manages the infrastructure. Available options:

- create\_before\_destroy
- prevent\_destroy
- ignore\_changes

<https://www.terraform.io/docs/configuration/resources.html#lifecycle-lifecycle-customizations>

### **terraform init**

- Downloads and initializes the provider
- Should be used to initialize remote backend to store state files
- To upgrade to latest acceptable version of provider: terraform init - upgrade

### **terraform plan**

- Would display the changes to be made to the infra. Performs terraform refresh automatically.
- Can save the generated terraform plan to specific path
- Once saved the plan can then be used with terraform apply to be certain that only the saved plan can be applied

terraform plan -out=path

#### **Example:**

terraform plan -out=examplefile

terraform apply examplefile

### **terraform refresh**

Fetches the current state of the infrastructure. Updates state files based on infrastructure implementation

### **terraform apply**

- Applies the terraform code and creates resources in infrastructure
- Can concurrently apply the changes – default to 10 concurrent operations. Change default using –parallelism flag
- Apply command will ask user for approval before applying the changes. For auto approval, set the flag –auto-approve [terraform apply –auto-approve]
- If the state file is already saved to external file, then use –state=path flag to set the path to state file
- Can target the apply only to certain resource [terraform apply –target=aws\_instance.MYEC2]
- To restrict asking user to input for variables if default value is not specified for variable use –input=false [Environment variable alternative – export TF\_INPUT="false" or export TF\_INPUT=0]

### **Terraform destroy:**

**terraform destroy:** Destroys all resources

**terraform destroy -auto-approve:** Will not ask for confirmation from the user before deleting

**terraform destroy -target aws\_instance.MyEC2:** To delete only the specific resource

### **terraform show**

Displays the current terraform state (tfstate file info)

### **terraform fmt**

To format the source code files according to terraform standards

### **terraform console**

Opens up terraform console for trying out terraform functions

### **terraform validate**

Checks whether the configuration file is syntactically valid

### **Desired state vs Current state:**

**Desired state:** Infrastructure configuration defined in terraform

**Current state:** What exactly is currently running in infrastructure

Terraform will plan to match the desired state to the current state. If there is any difference between both, the desired state will take the preference.

### **Variable assignments:**

Four ways we can pass variable value:

#### Variable defaults:

```
variable "instance_type" {  
    default = "t2.micro"  
}
```

#### Command line flags:

```
terraform apply -var="instance_type=t2.small"
```

#### From a file:

- Create a file **terraform.tfvars** (by default terraform reads variable values from this file if it exists. Recommended for production use)

```
instance_type = "t2.large"
```

- If we want to create custom file to hold variable values ex: custom.tfvars, then we need to explicitly tell terraform to use the custom file during terraform plan/apply

```
terraform plan -var-file="custom.tfvars"
```

Environment variables: (below examples for linux)

```
export TF_VAR_<variable name> = "<value>"
```

```
export TF_VAR_instnace_type = "t2.large"
```

If no default value specified, terraform apply will prompt user to keyin the value at run time

### Data Types:

- **string**: Example: "t2.micro"

```
variable "instance_type" {  
    type = string  
    default = "t2.micro"  
}
```

```
Usage: resource "aws_instance" "MyEC2" { instance_type = var.instance_type }
```

**\* older way (before 0.12 verison) of referencing variable value: `${var.instance_type}` but still being used in certain circumstances**

- **number**: Example: 8080

```
variable "port_number" {  
    type = number  
    default = 8080  
}
```

```
Usage: resource "aws_security_group" "MySG" { ingress  
{ to_port = var.port_number }}
```

- **bool**: Example: true/false

```
variable "isDefault" {  
    type = bool  
    default = true  
}
```

```
Usage: resource "aws_vpc" "defaultVPC" { default = var.isDefault }
```

- **list**: Example: ["ap-southeast-1","ap-southeast-2"]

```
variable "list_regions" {  
    type = list(string)  
    default = ["ap-southeast-1","ap-southeast-2"]  
}
```

```
Usage: resource "aws_instance" "MyEC2" { instance_type = var.list_regions["ap-southeast-1"] }
```

- **map**: Key value pairs.

```
variable "instance_type_in_regions" {
```

```

    type = map
    default = {
        us-east-1 = "t2.micro"
        us-east-2 = "t2.micro"
        us-east-3 = "t2.micro"
    }
}
Usage: resource "aws_instance" "MyEC2"
{ instance_type = var.instance_type_in_regions ["us-east-1"]}

```

- **set**: Similar to list. List can have duplicate data. Set cannot have duplicate data.

```

variable "list_regions" {
    type = set(string)
    default = ["ap-southeast-1", "ap-southeast-2"]
}
Usage: resource "aws_instance" "MyEC2" { instance_type = var.list_regions["ap-southeast-1"]}

```

- **object**:

```

variable "docker_config" {
    type = list(object({
        internal = number
        external = number
        protocol = string
    }))
    default = [{
        internal = 8080
        external = 80
        protocol = "HTTP"
    }]
}

```

- **tuple**: Similar to list, but can have different data types in it

```

variable "db_params" {
    type = tuple([string, number, bool])
    default = ["ap-southeast-1", 3306, true]
}

```

**Count: Parameter**: Allows to scale the resource based on the counter parameter specified.

```

resource "aws_iam_user" "lbuser" {
    name = "loadbalancer"
    path = "/system"
    count = 3 # Creates 3 users with same name
}

```

### Count Index:

```
resource "aws_iam_user" "lbuser" {
  name = "loadbalancer.${count.index}" # Assigns a different name based
  on count index
  path = "/system"
  count = 3 # Creates 3 users with same name
}
```

### Can also work with list variable:

```
variable "elb_names" {
  type = list
  default = ["dev-lb", "stage-lb", "prod-lb"]
}
resource "aws_iam_user" "lbuser" {
  name = var.elb_names[count.index] # Picks a name from list based on
  count index
  path = "/system"
  count = 3
}
```

### Conditional Expression (Ternary operator):

- Shortcut for if-else condition.

```
variable "isDev" {
  default = true
}
resource "aws_instance" "MyEC2Dev" {
  ami = "test"
  instance_type = "t2.micro"
  count = var.isDev == true ? 2 : 0 # Create production instances only
  if isDev is true
}
resource "aws_instance" "MyEC2Prod" {
  ami = "test"
  instance_type = "t2.large"
  count = var.isDev == false ? 2 : 0 # Create production instances only
  if isDev is false
}
```

### Local Values:

A local value assign a name to an expression which could be used multiple times *in a file* without repeating it.

```

locals {
    common_tags = {
        Name = "Dev Instance"
        Team = "SharePoint"
    }
}
resource "aws_instance" "MyEC2Dev" {
    ami = "test"
    instance_type = "t2.micro"
    tags = local.common_tags
}
resource "aws_instance" "MyEC2Dev1" {
    ami = "test"
    instance_type = "t2.micro"
    tags = local.common_tags
}

```

### Functions:

<https://www.terraform.io/docs/configuration/functions.html>

### Data Sources:

Retrieves data from provider to be used in terraform configuration. Example: AMI id in different regions is different / Retrieve subnet id's from default VPC

```

# get default subnet
data "aws_vpc" "default" {
    default = true
}

# get subnet id's from default VPC
data "aws_subnet_ids" "default" {
    vpc_id = data.aws_vpc.default.id
}

```

### Debugging Terraform:

- Detailed logs can be enabled by setting TF\_LOG environment variable.
- Can be set to one of the log levels – TRACE, DEBUG, INFO, WARN or ERROR
- Logs can be exported to files by setting TF\_LOG\_PATH. (export TF\_LOG\_PATH=/tmp/tflogs.log)

### Loading of files in the folder:

- If there are 4 files – web.tf, app.tf, sg.tf etc, terraform will load all the configuration files within the directory in alphabetical order
- Instead of writing all the configuration in one file, we can split into multiple files



## Dynamic Blocks:

- There could be cases where we need to add repeatable nested blocks of code which when implemented leads to repeatable and long code which is unmanageable

### Example AWS Security Group implementation which can have multiple ingress rules:

```
resource "aws_security_group" "SG" {
  ingress {
    from_port = 8080
    to_port   = 8080
    protocol = "HTTP"
    cidr_blocks = ["0.0.0.0/0"]
  }
  ingress {
    from_port = 8081
    to_port   = 8081
    protocol = "HTTP"
    cidr_blocks = ["0.0.0.0/0"]
  }
}
```

- Dynamic Blocks can help to dynamically construct repeatable code ex: "ingress" rule above.
- Can be provided inside resource, data, provider and provisioner blocks

### Example with dynamic block:

```
variable ingress_ports {
  type = list(number)
  default = [8080,8081]
}

dynamic "ingress" {
  for_each = var.ingress_ports
  iterator = port # optional argument, if omitted then use the label of the dynamic block (ingress). Otherwise use the iterator name in below code instead of "ingress" (ex: port.value)
  content {
    from_port = ingress.value
    to_port   = ingress.value
    protocol = "HTTP"
    cidr_blocks = ["0.0.0.0/0"]
  }
}
```

### Tainting Resources:

- Allows terraform-managed resource as tainted, forcing it to be destroyed and recreated on next apply.
- Updates the state file (tfstate file) for the resource to "tainted"
  - Will not modify the infrastructure
  - Will affect the resources that depend on the tainted resource

```
terraform taint <resource name>
```

### Example:

```
resource "aws_instance" "MyEC2" {  
    ami = "test"  
    instance_type = "t2.micro"  
}
```

```
terraform taint aws_instance.MyEC2
```

```
terraform apply (to destroy existing resource and create new resource)
```

### Splat Expressions and Output values:

- A splat expression provides a more concise way to express a common operation that could otherwise be performed with a for expression.

```
resource "aws_instance" "MyEC2Dev" {  
    ami = "test"  
    instance_type = "t2.micro"  
    count = 3  
}
```

```
# output value to output value of resource
```

```
output "arns"{  
    # [*] will get the arn's for all 3 instances created  
    value = aws_instance. MyEC2Dev[*].arn  
  
    description = "description of the output variable"  
  
    # By default false. True will prevent field's value to be displayed  
    # during output. However, the state file still shown un-encrypted value  
    sensitive = true  
}
```

### Terraform Graph:

- Provides graphical representation of the configuration.
- Output is in DOT format which can be converted to image
- DOT file can be visualized using Graphviz tool

```
terraform graph > graph.dot
```

### Provisioners:

Provisioners are used to execute scripts on a local or remote machine as part of resource creation or destruction.

```
resource "aws_instance" "MyEC2Dev" {
  ami = "test"
  instance_type = "t2.micro"

  provisioner "remote-exec" {
    # What scripts to run during startup
    inline {
      "sudo amazon-linux-extras install -y nginx1.12",
      "sudo systemctl start nginx"
    }

    # How to make connection to the VM instance
    connection {
      type = "ssh"
      host = self.public_ip
      user = "ec2-user"
      private_key = "${file("./terraform.pem")}"
    }
  }
}
```

### Types:

1. **local-exec**: Invokes local executable after resource is created. No "connection" block is required as the provisioner executes the command locally
2. **remote-exec**: Invokes a script on a remote resource after it is created. This can be used to run a configuration management tool like Ansible, Puppet, Chef etc. Supports both ssh and winrm connections

### Modules:

- Provides reusable code. Avoids repetitive code
- External modules can be referenced using "module" block

#### **Example:** EC2.tf

```
resource "aws_instance" "MyEC2Dev" {
  ami = "test"
  instance_type = "t2.micro"
}
```

#### Main.tf

```
module "myec2" {
```

```
        Source = "./EC2"
    }
```

In the above example, since "t2.micro" is hardcoded, the code importing the module cannot update the instance\_type.

If the instance\_type needs to be overwritten in the calling code, the instance\_type defined in module should be getting its value from a variable.

#### Example:

```
variable "instance_type"{
    default = "t2.micro"
}
resource "aws_instance" "MyEC2Dev" {
    ami = "test"
    instance_type = var.instance_type
}
```

With the above example, the calling code can override the instance\_type value.

#### Main.tf

```
module "myec2" {
    Source = "./EC2"
    Instance_type = "t2.large"
}
```

### **Terraform Registry:**

Pre-built modules that allows for quickly deploying common infrastructure configurations.

<https://registry.terraform.io/>

```
module "security-group" {
    source = "terraform-aws-modules/security-group/aws"
    version = "3.12.0"
    # insert the 2 required variables here
}
```

The above module can be used in our existing terraform configuration to create security groups.

### **Workspaces:**

Allows for separation of environments.

### **Commands:**

terraform workspace show: Show current workspace

terraform workspace list: To list all existing workspaces

terraform workspace select <workspace name>: To switch to the workspace specified

terraform workspace delete <workspace name>: To delete the workspace specified

terraform workspace new <workspace name>: To create new workspace specified

- In code, use terraform.workspace to retrieve the current workspace name
- Terraform.tfstate.d: Folder created by terraform to maintain **custom** workspaces. Each custom workspace (folder) contains terraform.tfstate file when terraform apply is ran.
- The terraform.tfstate file in root folder applies to "default" workspace

## State Management:

### Git:

- For team collaboration, state file can be stored in VCS like Git
- Problem: Sensitive info is stored in terraform.tfstate file which should be avoided

### Remote Backend:

Two types of backend:

- Standard Backend Type: Used for state storage and locking
- Enhanced Backend Type: Standard Backend Features + Remote Management

### Example:

```
terraform{
    backend "s3" {
        bucket = "bucket name"
        key = "key"
        region = "us-east-1"
    }
}
```

### State Locking:

- During performing any action, terraform would lock the state file
- When the state file is stored in remote backend, state locking is not available
- For S3, we can use DynamoDB (NoSQL db) for record locking

### Terraform commands for state management:

terraform state list: Lists all resources from state file

terraform state mv: Moves item with terraform state. Useful when renaming an existing resource without destroying and recreating it. Takes a backup of state prior to saving any changes.

terraform state pull: Manually download and output the state from remote state

terraform state push: Manually upload a local state file to remote state  
terraform state rm: Remove items from terraform state. Items removed from terraform state are not monitored. Resources created in the infra are not removed  
terraform state show: Show the attributes of a single resource in the state

<to be updated>

### **Terraform Import:**

Links manually created resource with the manually created terraform configuration

Example: terraform import aws\_instance.EC2 <instance id taken from AWS>

### **Terraform Cloud:**

<to be updated>